

Julia: a natural language for computational biology

Joe Greener
Research Associate, UCL
jgreener64.github.io



Talk outline

- Computational biology
- Julia
- BioJulia
- BioStructures.jl
- Bio3DView.jl
- Molly.jl
 - Simulating an ideal gas
 - Simulating diatomic molecules
 - Simulating proteins
- Differentiable molecular simulation
- Assessing Julia
- The future



What is biology?

“Biology is the natural science that studies life and living organisms, including their physical structure, chemical processes, molecular interactions, physiological mechanisms, development and evolution.” - Wikipedia




From Manning et al,
Science, 2002

What *is* biology?

- Complex systems
- Multi-scale systems
- Lots of data
- Incomplete, noisy data
- Rules have exceptions

Computational biology addresses this by developing and applying data analysis, mathematical modeling and simulation techniques to biology.



Protein structure prediction

- My day job involves predicting the structure of proteins from the sequence of amino acids that makes them up
- We do this using deep learning (CNNs/ResNets) on evolutionary information
- <https://github.com/psipred/DMPfold>

```
MENFQKVEKIGEGTYGVVYKARNKLTGEVVALKKIRLDTETEGVPSTAIRESLLKE  
LNHPNIVKLLDVIHTENKLYLVFEFLHQDLKKFMDASALTGIPLPLIKSYLFQLLQG  
LAFCHSHRVLHRDLKPQNLLINTEGAIKLADFLARAFGVPVRTYTHEVVTWYRAP  
EILLGCKYYSTAVDIWSLGCIFAEMVTRRALFPGDSEIDQLFRIFRTLGTPEVWWP  
GVTSMPTYKPSFPKWARQDFSKVVPPLDEDGRSLLSQMLHYDPNKRISAKAALHPF  
FQDVTKPVPHRLR
```



Julia

- High-level programming language
- Fast → addresses the two-language problem
- Free, open source
- Strong community
- Dynamically typed
- Multiple dispatch



```
[6]: abstract type Animal end

      struct Dog <: Animal
          name::String
      end

      struct Cat <: Animal
          name::String
          secret_name::String
      end

      sayname(a::Cat) = println("Maiow ", a.name)

      sayname(a::Dog) = println("Woof ", a.name)

      function meet(a::Animal, b::Animal)
          println(a.name, " meets ", b.name)
      end

      dog = Dog("Buster")
      cat = Cat("Salem", "??")

      sayname(dog)
```

Woof Buster

```
[7]: sayname(cat)
```

Maiow Salem

```
[8]: meet(dog, cat)
```

Buster meets Salem

BioJulia

- Fast, open, easy, software for biology
- The bioinformatics infrastructure for the Julia language



Efficient sequence type

```
In [1]: using BioSequences
```

```
In [2]: d = DNASequence("TTANC")
```

```
Out[2]: 5nt DNA Sequence:  
TTANC
```

```
In [3]: complement(d)
```

```
Out[3]: 5nt DNA Sequence:  
AATNG
```

BioJulia

File parsers

```
In [4]: reader = FASTA.Reader(open("myseq.fa", "r"))
        for record in reader
            # Do something
        end
        close(reader)
```

Alignments


```
In [10]: using BioAlignments
```

```
In [11]: s1 = dna"CCTAGGAGGG"
        s2 = dna"ACCTGGTATGATAGCG"
        scoremodel = AffineGapScoreModel(EDNAFULL, gap_open=-5, gap_extend=-1)
        res = pairalign(GlobalAlignment(), s1, s2, scoremodel)
```

```
Out[11]: PairwiseAlignmentResult{Int64,BioSequence{DNAAlphabet{4}},BioSequence{
DNAAlphabet{4}}}:
  score: 13
  seq:  0 -CCTAGG-----AGGG 10
        ||| ||         || |
  ref:  1 ACCT-GGTATGATAGCG 16
```


BioStructures.jl

- Read, write and manipulate macromolecular structures
- <https://github.com/BioJulia/BioStructures.jl>



PDB format

ATOM	1	N	MET	A	1	26.981	53.977	40.085	1.00	40.83	N
ATOM	2	CA	MET	A	1	26.091	52.849	39.889	1.00	37.14	C
ATOM	3	C	MET	A	1	26.679	52.163	38.675	1.00	30.15	C
ATOM	4	O	MET	A	1	27.020	52.865	37.715	1.00	27.59	O
ATOM	5	CB	MET	A	1	24.677	53.310	39.580	1.00	38.06	C
ATOM	6	CG	MET	A	1	23.624	52.189	39.442	1.00	46.67	C
ATOM	7	SD	MET	A	1	21.917	52.816	39.301	1.00	61.54	S
ATOM	8	CE	MET	A	1	21.930	53.926	37.910	1.00	51.17	C
ATOM	9	N	ARG	A	2	26.861	50.841	38.803	1.00	28.23	N
ATOM	10	CA	ARG	A	2	27.437	49.969	37.786	1.00	25.76	C

mmCIF format

ATOM	1	N	N	.	MET	A	1	1	?	26.981	53.977	40.085	1.00	40.83	?	1	MET	A	N	1
ATOM	2	C	CA	.	MET	A	1	1	?	26.091	52.849	39.889	1.00	37.14	?	1	MET	A	CA	1
ATOM	3	C	C	.	MET	A	1	1	?	26.679	52.163	38.675	1.00	30.15	?	1	MET	A	C	1
ATOM	4	O	O	.	MET	A	1	1	?	27.020	52.865	37.715	1.00	27.59	?	1	MET	A	O	1
ATOM	5	C	CB	.	MET	A	1	1	?	24.677	53.310	39.580	1.00	38.06	?	1	MET	A	CB	1
ATOM	6	C	CG	.	MET	A	1	1	?	23.624	52.189	39.442	1.00	46.67	?	1	MET	A	CG	1
ATOM	7	S	SD	.	MET	A	1	1	?	21.917	52.816	39.301	1.00	61.54	?	1	MET	A	SD	1
ATOM	8	C	CE	.	MET	A	1	1	?	21.930	53.926	37.910	1.00	51.17	?	1	MET	A	CE	1
ATOM	9	N	N	.	ARG	A	1	2	?	26.861	50.841	38.803	1.00	28.23	?	2	ARG	A	N	1
ATOM	10	C	CA	.	ARG	A	1	2	?	27.437	49.969	37.786	1.00	25.76	?	2	ARG	A	CA	1

BioStructures.jl

- Read, write and manipulate macromolecular structures

```
In [25]: using BioStructures
```

```
In [26]: struc = read("1AKE.pdb", PDB)
```

```
Out[26]: ProteinStructure 1AKE.pdb with 1 models, 2 chains (A,B), 428 residues, 3804 atoms
```

```
In [27]: struc[1]["A"]
```

```
Out[27]: Chain A with 214 residues, 242 other molecules, 1954 atoms
```

```
In [28]: countresidues(struc[1]["A"], standardselector)
```

```
Out[28]: 214
```

```
In [29]: for at in collectatoms(struc, calphaselector)[1:5]
          println(atomname(at), " ", coords(at))
        end
```

```
CA [26.091, 52.849, 39.889]
CA [27.437, 49.969, 37.786]
CA [24.961, 47.988, 35.671]
CA [25.194, 44.925, 33.36]
CA [22.428, 44.503, 30.712]
```

BioStructures.jl

- Spatial calculations

```
In [39]: # Print the PDB records for all Cα atoms within 4 Å of residue 38
for at in collectatoms(struc['A'], calphaselector)
    if distance(struc['A'][38], at) < 4.0 && resnumber(at) != 38
        println(pdbline(at))
    end
end
```

ATOM	270	CA	ALA	A	37	33.778	51.895	15.373	1.00	27.13	C
ATOM	280	CA	VAL	A	39	36.426	48.279	12.266	1.00	21.10	C
ATOM	296	CA	SER	A	41	38.507	53.491	11.905	1.00	33.63	C
ATOM	302	CA	GLY	A	42	40.955	50.617	11.185	1.00	29.51	C
ATOM	329	CA	GLY	A	46	39.438	45.942	17.303	1.00	17.17	C
ATOM	333	CA	LYS	A	47	39.679	46.226	13.467	1.00	25.84	C

BioStructures.jl


- Spatial calculations

```
In [33]: dists = DistanceMap(collectatoms(struc['A'], cbetaselector))
```

```
Out[33]: Distance map of size (214, 214)
```

```
In [34]: using Plots  
plot(dists)
```

```
Out[34]:
```




BioStructures.jl

- Access the Protein Data Bank (PDB)

```
In [ ]: # Calculate the cysteine fraction of every structure in the PDB
l = pdbentrylist()
for p in l
    downloadpdb(p, file format=MMCIF) do fp
        s = read(fp, MMCIF)
        nres = countresidues(s, standardselector)
        if nres > 0
            frac = countresidues(s, standardselector, x -> resname(x) == "CYS") / nres
            println(p, " ", round(frac, digits=2))
        end
    end
end
```

Bio3DView.jl

- Visualisation of macromolecular structures
- Wrapper round 3Dmol.js (Rego and Koes, *Bioinformatics*, 2015)
- Works in IJulia or Blink popup window
- <https://github.com/jgreener64/Bio3DView.jl>




Molly.jl

- Molecular dynamics (MD) is a computational technique used to explore the movement and interaction of molecules
- Molly.jl is a proof-of-concept for MD in Julia
- <https://github.com/jgreener64/Molly.jl>

GROMACS - Project Cost	
Include	Avg. Salary
Markup And Code ▾	\$ 55000 /year
Codebase	Effort (est.)
814,364 Lines	221 Person Years
Estimated Cost	\$12,140,504
Updated Jun 24, 2019 more at Open Hub	

From gromacs.org



From Jung et al, *J Comp Chem*, 2019

Simulating an ideal gas

- Create some atoms with the relevant parameters defined

```
In [1]: using Molly
```

n atoms = 100

```
mass = 10.0
```

```
atoms = [Atom(mass=mass,  $\sigma$ =0.3,  $\epsilon$ =0.2) for i in 1:n atoms]
```

Out[1]: 100-element Array{Atom,1}:

```
Atom("", "", 0, "", 0.0, 10.0, 0.3, 0.2)
```

```
Atom("", "", 0, "", 0.0, 10.0, 0.3, 0.2)
```

```
Atom("", "", 0, "", 0.0, 10.0, 0.3, 0.2)
```

```
Atom("", "", 0, "", 0.0, 10.0, 0.3, 0.2)
```

```
Atom("", "", 0, "", 0.0, 10.0, 0.3, 0.2)
```

```
Atom("", "", 0, "", 0.0, 10.0, 0.3, 0.2)
```

```
Atom(" ", " ", 0, " ", 0.0, 10.0, 0.3, 0.2)
```

```
Atom(" ", " ", 0, " ", 0.0, 10.0, 0.3, 0.2)
```

```
Atom(" ", " ", 0, " ", 0.0, 10.0, 0.3, 0.2)
```

```
Atom(" ", " ", 0, " ", 0.0, 10.0, 0.3, 0.2)
```

Atom("": "", 0, "", 0.0, 10.0, 0.3, 0.2)

```
Atom(" ", " ", 0, " ", 0.0, 10.0, 0.3, 0.2)
```

```
Atom(" ", " ", 0, " ", 0.0, 10.0, 0.3, 0.2)
```

•

Simulating an ideal gas

- Define some starting coordinates and velocities

```
In [2]: box_size = 2.0 # nm  
coords = [Coordinates(rand(3) .* box_size) for i in 1:n_atoms]  
  
temperature = 298 # K  
velocities = [Velocity(mass, temperature) for i in 1:n_atoms]
```

```
Out[2]: 100-element Array{Velocity,1}:  
[-3.62082, -4.11298, -1.80747]  
[-6.06135, 2.48857, -0.968609]  
[3.5725, -1.64122, 2.61146]  
[-8.20442, -1.96685, 12.5213]  
[6.7406, 8.06838, 2.47006]  
[6.26042, -5.08734, 7.4645]  
[-5.41582, 7.2642, -0.0616184]  
[10.4321, -2.95435, -1.25541]  
[4.51567, 6.09565, 4.18389]  
[-5.45313, 1.96933, -1.17946]  
[5.50508, -5.97522, -2.35687]  
[-0.969843, 2.73667, 4.7403]  
[-2.38129, 4.79214, -6.1126]  
⋮
```

Simulating an ideal gas

- Now we can define our dictionary of general interactions, i.e. those between most or all atoms
- Because we have defined the relevant parameters for the atoms, we can use the built-in Lennard Jones type

```
In [3]: general_inters = Dict("LJ" => LennardJones())
```

```
Out[3]: Dict{String,LennardJones} with 1 entry:  
        "LJ" => LennardJones(false)
```

Simulating an ideal gas

- Define and run the simulation
- Use an Andersen thermostat to keep a constant temperature
- Log the temperature and coordinates every 100 steps

```
In [4]: s = Simulation(  
    simulator=VelocityVerlet(), # Use velocity Verlet integration  
    atoms=atoms,  
    general_inters=general_inters,  
    coords=coords,  
    velocities=velocities,  
    temperature=temperature,  
    box_size=box_size,  
    thermostat=AndersenThermostat(1.0), # Coupling constant of 1.0  
    loggers=[TemperatureLogger(100), CoordinateLogger(100)],  
    timestep=0.002, # ps  
    n_steps=100_000  
)  
  
simulate!(s)
```

Starting simulation

Progress: 99% |  | ETA: 0:00:00

Simulating an ideal gas

- Plot the simulation

```
In [26]: using Plots
          pyplot(leg=false)

          coords = s.loggers[2].coords
          temps = s.loggers[1].temperatures

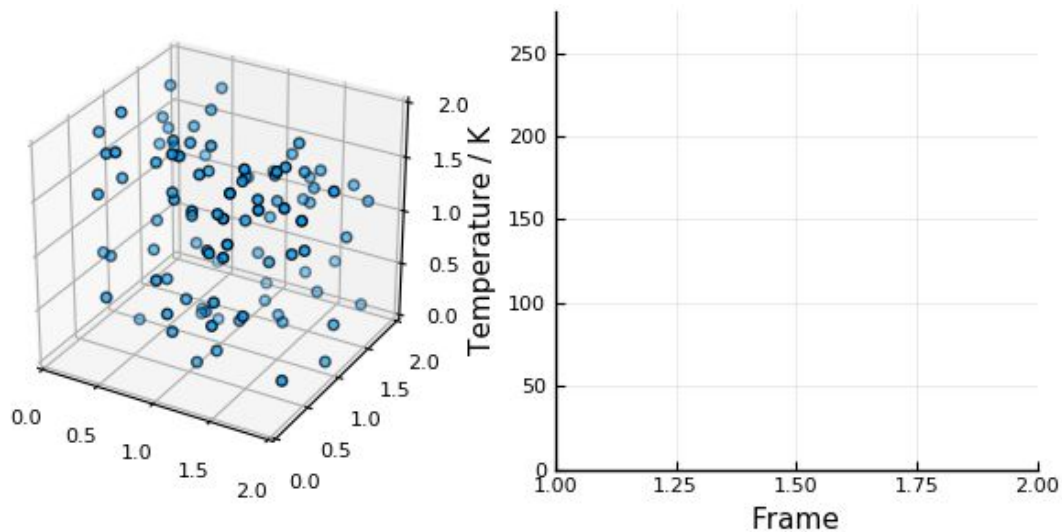
          splitcoords(coord) = [c[1] for c in coord], [c[2] for c in coord], [c[3] for c in coord]

          @gif for (i, coord) in enumerate(coords)
              l = @layout [a b{0.7h}]

              cx, cy, cz = splitcoords(coord)
              p = scatter(cx, cy, cz,
                          xlims=(0, box_size),
                          ylims=(0, box_size),
                          zlims=(0, box_size),
                          layout=l
                        )

              plot!(p[2],
                    temps[1:i],
                    xlabel="Frame",
                    ylabel="Temperature / K",
                    xlims=(1, i),
                    ylims=(0.0, maximum(temps[1:i]))
                  )
          end
```

Simulating an ideal gas



Simulating a diatomic molecule

- We can define specific interactions between groups of atoms, e.g. bonds
- Use the built-in bond type to place a harmonic constraint between paired atoms

```
In [13]: bonds = [Bond((i * 2) - 1, i * 2, 0.1, 300_000) for i in 1:(n_atoms / 2)]  
specific_inter_lists = Dict("Bonds" => bonds)
```


```
In [12]: coords = Coordinates[]  
for i in 1:(n_atoms / 2)  
    c = rand(3) .* box_size  
    push!(coords, Coordinates(c))  
    push!(coords, Coordinates(c + [0.1, 0.0, 0.0]))  
end  
  
velocities = [Velocity(mass, temperature) for i in 1:n_atoms]
```

Simulating a diatomic molecule

- Use a neighbour list to speed up the Lennard Jones calculation
- We will use the provided neighbour finder but you could write your own

```
In [15]: neighbour_finder = DistanceNeighbourFinder(trues(n_atoms, n_atoms), 10, 1.2)
```

Simulating a diatomic molecule



Simulating a protein

- In addition to Lennard Jones forces and bonds, proteins also have angles, dihedral angles (torsions) and electrostatics
- Molly.jl has basic functionality to read in topology and forcefield data from GROMACS format files

```
In [27]: timestep = 0.0002 # ps
          temperature = 298 # K
          n_steps = 5000


          atoms, specific_inter_lists, general_inters, nb_matrix, coords, box_size = readinputs(
              joinpath(dirname(pathof(Molly)), "..", "data", "5XER", "gmx_top_ff.top"),
              joinpath(dirname(pathof(Molly)), "..", "data", "5XER", "gmx_coords.gro"))
```

Simulating a protein

- Simulate as before

```
In [*]: s = Simulation(  
    simulator=VelocityVerlet(),  
    atoms=atoms,  
    specific_inter_lists=specific_inter_lists,  
    general_inters=general_inters,  
    coords=coords,  
    velocities=[Velocity(a.mass, temperature) for a in atoms],  
    temperature=temperature,  
    box_size=box_size,  
    neighbour_finder=DistanceNeighbourFinder(nb_matrix, 10),  
    thermostat=AndersenThermostat(1.0),  
    loggers=[TemperatureLogger(10), StructureWriter(10, "traj_5XER_1ps.pdb")],  
    timestep=timestep,  
    n_steps=n_steps  
)  
  
simulate!(s)
```


Starting simulation

Progress: 5%

| ETA: 0:11:04

Simulating a protein

- 1 ps simulation viewed with VMD



Defining a force

- The entire definition for the electrostatic (Coulomb) force

```
In [ ]: struct Coulomb <: GeneralInteraction
        nl_only::Bool
      end
```

```
In [ ]: @fastmath @inbounds function update_accelerations!(accels::Vector{Acceleration},
                                                           inter::Coulomb,
                                                           s::Simulation,
                                                           i::Integer,
                                                           j::Integer)
    dx = vector1D(s.coords[i].x, s.coords[j].x, s.box_size)
    dy = vector1D(s.coords[i].y, s.coords[j].y, s.box_size)
    dz = vector1D(s.coords[i].z, s.coords[j].z, s.box_size)
    r2 = dx * dx + dy * dy + dz * dz
    if r2 > sqdist_cutoff_nb
        return
    end
    f = (coulomb_const * s.atoms[i].charge * s.atoms[j].charge) / sqrt(r2 ^ 3)
    accels[i].x += -f * dx
    accels[i].y += -f * dy
    accels[i].z += -f * dz
    accels[j].x += f * dx
    accels[j].y += f * dy
    accels[j].z += f * dz
  end
```

Differentiable molecular simulation


- Recent developments with Flux.jl make it possible to run molecular simulations and obtain gradients of parameters, e.g. forcefield parameters

```
[66]: b0learn = param(1.0)
      b0true = 4.0
      k = 0.01
      nsteps = 500

      function simulate(nsteps)
        coords = [3.0, 5.0]
        coords_last = copy(coords) + randn(2) * 0.01
        for i in 1:nsteps
          dist = abs(coords[2] - coords[1])
          force = k * (dist - b0learn) * abs(dist - b0learn)
          dir = coords[1] < coords[2] ? 1.0 : -1.0
          coords_next = 2 * coords - coords_last + [dir * force, -dir * force]
          coords = coords_next
          coords_last = coords
        end
        return coords
      end

      simulate(nsteps)
```

Step 1 / 500



Differentiable molecular simulation

```
[74]: function loss(b0true, nsteps)
      coords = simulate(nsteps)
      dist = abs(coords[2] - coords[1])
      return (dist - b0true) ^ 2
    end


    gs = Tracker.gradient(() -> loss(b0true, nsteps), params(b0learn))
    gs[b0learn]
```

```
[74]: -5.773390700589583 (tracked)
```

```
[75]: for i in 1:500
      gs = Tracker.gradient(() -> loss(b0true, nsteps), params(b0learn))
      update!(b0learn, -0.01 * gs[b0learn])
      i % 50 == 0 && println("Epoch ", i, " - b0learn ", b0learn)
    end
```

```
Epoch 50 - b0learn 2.6801969855276746 (tracked)
Epoch 100 - b0learn 3.574286748329535 (tracked)
Epoch 150 - b0learn 3.9043047080232394 (tracked)
Epoch 200 - b0learn 4.025176930121734 (tracked)
Epoch 250 - b0learn 4.069391236713611 (tracked)
Epoch 300 - b0learn 4.0855609902908805 (tracked)
Epoch 350 - b0learn 4.09147724816171 (tracked)
Epoch 400 - b0learn 4.093635443806043 (tracked)
Epoch 450 - b0learn 4.094422175776826 (tracked)
Epoch 500 - b0learn 4.094708733260627 (tracked)
```

Step 1 / 500



Julia for computational biology

Pros

- High-level language with nice syntax
- Fast
- Prototype and code in the same language
- Sane design
- Great community
- Making packages is relatively easy
- There is a rabbit hole to go down


Cons

- Compile time requires a workflow change
- Performance gotchas
- Effort required to shift from OOP
- Plotting is still getting there
- Libraries still being implemented
- Sometimes you have to go down the rabbit hole



The future

- Popularity of the language increasing
- Increasingly used in teaching and research
- Stability of Julia v1.0 has made things easier for package developers
- Exciting developments in the language:
 - Improved multithreading
 - GPU/TPU programming
 - Deep learning/differentiable programming
 - Static compilation



Acknowledgements

- UCL Bioinformatics Group
- European Research Council
- Julia community
- BioJulia, in particular:
 - Ben Ward (GitHub BenJWard)
 - Kenta Sato (bicycle1885)
 - Joel S (joelselvaraj)

